

# r4ds-ch3

2025-05-04

## Ch 3 Data transformation

### 3.1 Introduction

#### 3.1.1 Prerequisites

```
library(nycflights13)
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.2      v tibble    3.2.1
## v lubridate  1.9.4      v tidyr     1.3.1
## v purrr      1.0.4
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

dplyr overwrites some base R functions, so when want to use them have to use their full names  
e.g. stats::filter() When need to be precise, use packagename::functionname()

#### 3.1.2 nycflights13

nycflights13::flights contains 336,776 flights that departed from NYC in 2013

```
flights
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
```

```
## 8 2013 1 1 557 600 -3 709 723
## 9 2013 1 1 557 600 -3 838 846
## 10 2013 1 1 558 600 -2 753 745
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

flights is a tibble, a special type of data frame used by the tidyverse. Most important difference is how they print, tibbles are designed for large datasets so they only show the first few rows and columns that can fit on the screen. In RStudio you can use `View(flights)` for an interactive view of all the data. You can also do `print(flights, width = Inf)` to show all columns.

```
glimpse(flights) # Or use glimpse()
```

```
## Rows: 336,776
## Columns: 19
## $ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2~
## $ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
## $ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
## $ dep_time  <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 558, ~
## $ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 600, ~
## $ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2, -1~
## $ arr_time  <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 849, ~
## $ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 851, ~
## $ arr_delay <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7, -1~
## $ carrier   <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6", "~
## $ flight    <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301, 4~
## $ tailnum   <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N394~
## $ origin    <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LGA", ~
## $ dest      <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IAD", ~
## $ air_time  <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149, 1~
## $ distance  <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 733, ~
## $ hour      <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6, 6~
## $ minute    <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 59, 0~
## $ time_hour <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-01 0~
```

Variable names are followed by abbreviations for the type of variable: integers, doubles (real numbers), character (strings), dtm for date-time.

### 3.1.3 dplyr basics

What dplyr verbs (functions) have in common 1. First argument is always a data frame 2. Subsequent arguments describe which columns to operate on using the variable names (without quotes) 3. Output is always a new data frame

Combine multiple verbs with the pipe `|>` Pipe takes thing on its left and passes it to function on its right `x |> f(y)` is equivalent to `f(x, y)` `x |> f(y) |> g(z)` is equivalent to `g(f(x, y), z)` Pronounce the pipe as “then”

```
flights |>
  filter(dest == "IAH") |>
  group_by(year, month, day) |>
```

```
summarize(
  arr_delay = mean(arr_delay, na.rm = TRUE)
)
```

## 'summarise()' has grouped output by 'year', 'month'. You can override using the  
## '.groups' argument.

```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day arr_delay
##   <int> <int> <int>    <dbl>
## 1  2013     1     1     17.8
## 2  2013     1     2      7
## 3  2013     1     3     18.3
## 4  2013     1     4     -3.2
## 5  2013     1     5     20.2
## 6  2013     1     6      9.28
## 7  2013     1     7     -7.74
## 8  2013     1     8      7.79
## 9  2013     1     9     18.1
## 10 2013     1    10      6.68
## # i 355 more rows
```

The code above groups flights whose destination is IAH and displays the mean arrival delay for each day  
dplyr verbs are organized into 4 groups based on what they operate on: rows, columns, groups, or tables

## 3.2 Rows

Most important for rows are 1. `filter()` which changes which rows are present without changing their order  
2. `arrange()` which changes the order of rows without changing which are present

### 3.2.1 filter()

Allows you to keep rows based on values of columns First argument is data frame, the next are conditions  
that must be true for the row to be kept

Find all flights that departed more than 2 hours late

```
flights |>
  filter(dep_delay > 120)
```

```
## # A tibble: 9,723 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>    <int>         <int>    <dbl>    <int>         <int>
## 1  2013     1     1      848           1835      853     1001           1950
## 2  2013     1     1      957           733      144     1056            853
## 3  2013     1     1     1114           900      134     1447           1222
## 4  2013     1     1     1540          1338      122     2020           1825
## 5  2013     1     1     1815          1325      290     2120           1542
## 6  2013     1     1     1842          1422      260     1958           1535
```

```
## 7 2013 1 1 1856 1645 131 2212 2005
## 8 2013 1 1 1934 1725 129 2126 1855
## 9 2013 1 1 1938 1703 155 2109 1823
## 10 2013 1 1 1942 1705 157 2124 1830
## # i 9,713 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Can combine conditions with & or , to indicate “and” and | to indicate “or”

```
# Flights that departed on Jan 1
flights |>
  filter(month == 1 & day == 1)
```

```
## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     1     1     517           515           2     830           819
## 2 2013     1     1     533           529           4     850           830
## 3 2013     1     1     542           540           2     923           850
## 4 2013     1     1     544           545          -1    1004          1022
## 5 2013     1     1     554           600          -6     812           837
## 6 2013     1     1     554           558          -4     740           728
## 7 2013     1     1     555           600          -5     913           854
## 8 2013     1     1     557           600          -3     709           723
## 9 2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # i 832 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Flights that departed in Jan or Feb
flights |>
  filter(month == 1 | month == 2)
```

```
## # A tibble: 51,955 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     1     1     517           515           2     830           819
## 2 2013     1     1     533           529           4     850           830
## 3 2013     1     1     542           540           2     923           850
## 4 2013     1     1     544           545          -1    1004          1022
## 5 2013     1     1     554           600          -6     812           837
## 6 2013     1     1     554           558          -4     740           728
## 7 2013     1     1     555           600          -5     913           854
## 8 2013     1     1     557           600          -3     709           723
## 9 2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # i 51,945 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Useful shortcut to combine | and == which is %in%

```
flights |>
  filter(month %in% c(1, 2))
```

```
## # A tibble: 51,955 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
##10  2013     1     1     558           600          -2     753           745
## # i 51,945 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Filter() creates a new data frame then prints it, doesn't modify original dplyr functions never modify their inputs To save the result, do so with the assignment operator

```
jan1 <- flights |>
  filter(month == 1 & day == 1)
```

### 3.2.2 Common mistakes

Using = instead of == to test for equality, filter() lets you know when this happens

```
#flights |>
#filter(month = 1)
```

Another mistake is to write “or” statements like you would in English

```
flights |>
  filter(month == 1 | 2)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
```

```
## 8 2013 1 1 557 600 -3 709 723
## 9 2013 1 1 557 600 -3 838 846
## 10 2013 1 1 558 600 -2 753 745
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

No error, but first checks condition month == 1 and then 2, 2 is always true so this doesn't filter anything

### 3.2.3 arrange()

Changes order of rows based on value of the columns Takes a data frame and set of column names to order by If provide more than 1 column, each successive column is used to break ties

Sort by departure time, earliest years first, then within a year the earliest months, etc.

```
flights |>
  arrange(year, month, day, dep_time)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     1     1     517           515           2     830           819
## 2 2013     1     1     533           529           4     850           830
## 3 2013     1     1     542           540           2     923           850
## 4 2013     1     1     544           545          -1    1004          1022
## 5 2013     1     1     554           600          -6     812           837
## 6 2013     1     1     554           558          -4     740           728
## 7 2013     1     1     555           600          -5     913           854
## 8 2013     1     1     557           600          -3     709           723
## 9 2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Can use desc() on a column inside of arrange() to re-order data frame based on that column in descending order

This orders flights from most to least delayed

```
flights |>
  arrange(desc(dep_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     1     9     641           900        1301    1242          1530
## 2 2013     6    15    1432          1935        1137    1607          2120
## 3 2013     1    10    1121          1635        1126    1239          1810
## 4 2013     9    20    1139          1845        1014    1457          2210
```

```
## 5 2013 7 22 845 1600 1005 1044 1815
## 6 2013 4 10 1100 1900 960 1342 2211
## 7 2013 3 17 2321 810 911 135 1020
## 8 2013 6 27 959 1900 899 1236 2226
## 9 2013 7 22 2257 759 898 121 1026
## 10 2013 12 5 756 1700 896 1058 2020
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Number of rows doesn't change, not filtering, only re-ordering

### 3.2.4 distinct()

Finds all unique rows in dataset Most of time will want distinct combo of some variables, so can also optionally supply column names

```
# Remove duplicate rows if there are any
flights |>
  distinct()
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     1     1     517           515           2     830           819
## 2 2013     1     1     533           529           4     850           830
## 3 2013     1     1     542           540           2     923           850
## 4 2013     1     1     544           545          -1    1004          1022
## 5 2013     1     1     554           600          -6     812           837
## 6 2013     1     1     554           558          -4     740           728
## 7 2013     1     1     555           600          -5     913           854
## 8 2013     1     1     557           600          -3     709           723
## 9 2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Find all unique origin and destination pairs
flights |>
  distinct(origin, dest)
```

```
## # A tibble: 224 x 2
##   origin dest
##   <chr> <chr>
## 1 EWR   IAH
## 2 LGA   IAH
## 3 JFK   MIA
## 4 JFK   BQN
## 5 LGA   ATL
```

```
## 6 EWR    ORD
## 7 EWR    FLL
## 8 LGA    IAD
## 9 JFK    MCO
## 10 LGA   ORD
## # i 214 more rows
```

If want to keep all other columns when filtering for unique rows, can use `.keep_all = TRUE`

```
flights |>
  distinct(origin, dest, .keep_all = TRUE)
```

```
## # A tibble: 224 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     517           515         2      830           819
## 2  2013     1     1     533           529         4      850           830
## 3  2013     1     1     542           540         2      923           850
## 4  2013     1     1     544           545        -1     1004          1022
## 5  2013     1     1     554           600        -6      812           837
## 6  2013     1     1     554           558        -4      740           728
## 7  2013     1     1     555           600        -5      913           854
## 8  2013     1     1     557           600        -3      709           723
## 9  2013     1     1     557           600        -3      838           846
## 10 2013     1     1     558           600        -2      753           745
## # i 214 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Not coincidence all from Jan 1, `distinct()` finds first occurrence of row in dataset and discards rest

To find number of occurrences, swap `distinct()` with `count()` With `sort = TRUE`, arranges them in descending order of number of occurrences

```
flights |>
  count(origin, dest, sort = TRUE)
```

```
## # A tibble: 224 x 3
##   origin dest      n
##   <chr> <chr> <int>
## 1 JFK    LAX    11262
## 2 LGA    ATL    10263
## 3 LGA    ORD     8857
## 4 JFK    SFO     8204
## 5 LGA    CLT     6168
## 6 EWR    ORD     6100
## 7 JFK    BOS     5898
## 8 LGA    MIA     5781
## 9 JFK    MCO     5464
## 10 EWR    BOS     5327
## # i 214 more rows
```



### 3.2.5 Exercises

1. In a single pipeline for each condition, find all flights that meet the condition:

a) Arrival delay of two or more hours

```
flights |>
  filter(arr_delay >= 120)
```

```
## # A tibble: 10,200 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     811             630          101    1047             830
## 2  2013     1     1     848             1835         853    1001             1950
## 3  2013     1     1     957             733          144    1056             853
## 4  2013     1     1    1114             900          134    1447             1222
## 5  2013     1     1    1505             1310         115    1638             1431
## 6  2013     1     1    1525             1340          105    1831             1626
## 7  2013     1     1    1549             1445           64    1912             1656
## 8  2013     1     1    1558             1359         119    1718             1515
## 9  2013     1     1    1732             1630           62    2028             1825
## 10 2013     1     1    1803             1620          103    2008             1750
## # i 10,190 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

b) Flew to Houston (IAH or HOU)

```
flights |>
  filter(dest %in% c("IAH", "HOU"))
```

```
## # A tibble: 9,313 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     623             627          -4     933             932
## 4  2013     1     1     728             732          -4    1041             1038
## 5  2013     1     1     739             739           0    1104             1038
## 6  2013     1     1     908             908           0    1228             1219
## 7  2013     1     1    1028             1026           2    1350             1339
## 8  2013     1     1    1044             1045          -1    1352             1351
## 9  2013     1     1    1114             900         134    1447             1222
## 10 2013     1     1    1205             1200           5    1503             1505
## # i 9,303 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

c) Were operated by United, American, or Delta

```
airlines # Look up airline codes
```

```
## # A tibble: 16 x 2
##   carrier name
##   <chr>   <chr>
## 1 9E      Endeavor Air Inc.
## 2 AA      American Airlines Inc.
## 3 AS      Alaska Airlines Inc.
## 4 B6      JetBlue Airways
## 5 DL      Delta Air Lines Inc.
## 6 EV      ExpressJet Airlines Inc.
## 7 F9      Frontier Airlines Inc.
## 8 FL      AirTran Airways Corporation
## 9 HA      Hawaiian Airlines Inc.
## 10 MQ     Envoy Air
## 11 OO     SkyWest Airlines Inc.
## 12 UA     United Air Lines Inc.
## 13 US     US Airways Inc.
## 14 VX     Virgin America
## 15 WN     Southwest Airlines Co.
## 16 YV     Mesa Airlines Inc.
```

```
flights |>
  filter(carrier %in% c("UA", "AA", "DL"))
```

```
## # A tibble: 139,504 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     554           600          -6     812           837
## 5  2013     1     1     554           558          -4     740           728
## 6  2013     1     1     558           600          -2     753           745
## 7  2013     1     1     558           600          -2     924           917
## 8  2013     1     1     558           600          -2     923           937
## 9  2013     1     1     559           600          -1     941           910
## 10 2013     1     1     559           600          -1     854           902
## # i 139,494 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

d) Departed in summer (July, August, and September)

```
flights |>
  filter(month %in% c(7:9))
```

```
## # A tibble: 86,326 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
```

```
## 1 2013 7 1 1 2029 212 236 2359
## 2 2013 7 1 2 2359 3 344 344
## 3 2013 7 1 29 2245 104 151 1
## 4 2013 7 1 43 2130 193 322 14
## 5 2013 7 1 44 2150 174 300 100
## 6 2013 7 1 46 2051 235 304 2358
## 7 2013 7 1 48 2001 287 308 2305
## 8 2013 7 1 58 2155 183 335 43
## 9 2013 7 1 100 2146 194 327 30
## 10 2013 7 1 100 2245 135 337 135
## # i 86,316 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

e) Arrived more than two hours late but didn't leave late

```
flights |>
  filter(dep_delay <= 0 & arr_delay > 120)
```

```
## # A tibble: 29 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1 2013     1    27    1419           1420        -1     1754           1550
## 2 2013    10     7    1350           1350         0     1736           1526
## 3 2013    10     7    1357           1359        -2     1858           1654
## 4 2013    10    16     657             700        -3     1258           1056
## 5 2013    11     1     658             700        -2     1329           1015
## 6 2013     3    18    1844           1847        -3         39           2219
## 7 2013     4    17    1635           1640        -5     2049           1845
## 8 2013     4    18     558             600        -2     1149           850
## 9 2013     4    18     655             700        -5     1213           950
## 10 2013     5    22    1827           1830        -3     2217           2010
## # i 19 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

f) Were delayed by at least an hour, but made up over 30 min in flight

```
flights |>
  filter(dep_delay >= 60 & (dep_delay - arr_delay) > 30)
```

```
## # A tibble: 1,844 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1 2013     1     1    2205           1720       285         46           2040
## 2 2013     1     1    2326           2130       116        131           18
## 3 2013     1     3    1503           1221       162       1803           1555
## 4 2013     1     3    1839           1700        99       2056           1950
## 5 2013     1     3    1850           1745        65       2148           2120
## 6 2013     1     3    1941           1759       102       2246           2139
```

```
## 7 2013 1 3 1950 1845 65 2228 2227
## 8 2013 1 3 2015 1915 60 2135 2111
## 9 2013 1 3 2257 2000 177 45 2224
## 10 2013 1 4 1917 1700 137 2135 1950
## # i 1,834 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

2. Sort flights to find the flights with the longest departure delays

```
flights |>
  arrange(desc(dep_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>      <int>         <int>
## 1 2013     1     9     641           900      1301      1242          1530
## 2 2013     6    15    1432          1935      1137      1607          2120
## 3 2013     1    10    1121          1635      1126      1239          1810
## 4 2013     9    20    1139          1845      1014      1457          2210
## 5 2013     7    22     845          1600      1005      1044          1815
## 6 2013     4    10    1100          1900       960      1342          2211
## 7 2013     3    17    2321           810       911       135          1020
## 8 2013     6    27     959          1900       899      1236          2226
## 9 2013     7    22    2257           759       898       121          1026
## 10 2013    12     5     756          1700       896      1058          2020
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Find the flights that left earliest in the morning

```
flights |>
  arrange(dep_time)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>      <int>         <int>
## 1 2013     1    13         1           2249       72       108          2357
## 2 2013     1    31         1           2100      181       124          2225
## 3 2013    11    13         1           2359        2       442          440
## 4 2013    12    16         1           2359        2       447          437
## 5 2013    12    20         1           2359        2       430          440
## 6 2013    12    26         1           2359        2       437          440
## 7 2013    12    30         1           2359        2       441          437
## 8 2013     2    11         1           2100      181       111          2225
## 9 2013     2    24         1           2245       76       121          2354
## 10 2013     3     8         1           2355        6       431          440
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

### 3. Sort flights to find the fastest flights

```
# First convert arrival and departure times to minutes
arrival <- function() {
  as.integer(substring(str_pad(flights$arr_time, 4, side = "left", pad = 0), 1, 2))*60 + as.integer(substring(str_pad(flights$arr_time, 4, side = "left", pad = 0), 3, 4))
}

departure <- function() {
  as.integer(substring(str_pad(flights$dep_time, 4, side = "left", pad = 0), 1, 2))*60 + as.integer(substring(str_pad(flights$dep_time, 4, side = "left", pad = 0), 3, 4))
}

# Subtract departure from arrival to determine complete flight time
flight_time <- function() {
  times <- arrival() - departure()
  # If negative, then add 1400 (means departed at night and arrived morning)
  times[times < 0 & !is.na(times)] <- times[times < 0 & !is.na(times)] + 1400
  times
}

# Arrange by flight time
flights |>
  arrange(flight_time())
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     6    12    2338             2129          129     17           2235
## 2  2013    12    29    2332             2155           97     14           2300
## 3  2013    11     6    2335             2215           80     18           2317
## 4  2013     2    25    2347             2145          122     30           2239
## 5  2013     8    13    2351             2152          119     35           2258
## 6  2013    10    11    2342             2030          192     27           2205
## 7  2013     2    26    2356             2000          236     41           2104
## 8  2013     1    24    2342             2159          103     28           2300
## 9  2013    12    23    2333             2155           98     19           2257
## 10 2013     3    10    2339             2200           99     25           2254
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

### 4. Was there a flight every day of 2013?

```
nrow(distinct(flights, month, day)) == 365
```

```
## [1] TRUE
```

Yes.

### 5. Which flights traveled the farthest distance?

```
flights |>
  arrange(desc(distance))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     857             900          -3    1516           1530
## 2  2013     1     2     909             900           9    1525           1530
## 3  2013     1     3     914             900          14    1504           1530
## 4  2013     1     4     900             900           0    1516           1530
## 5  2013     1     5     858             900          -2    1519           1530
## 6  2013     1     6    1019             900          79    1558           1530
## 7  2013     1     7    1042             900         102    1620           1530
## 8  2013     1     8     901             900           1    1504           1530
## 9  2013     1     9     641             900        1301    1242           1530
## 10 2013     1    10     859             900          -1    1449           1530
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Which traveled the least distance?

```
flights |>
  arrange(distance)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     7    27      NA             106          NA      NA           245
## 2  2013     1     3    2127             2129          -2    2222           2224
## 3  2013     1     4    1240             1200          40    1333           1306
## 4  2013     1     4    1829             1615         134    1937           1721
## 5  2013     1     4    2128             2129          -1    2218           2224
## 6  2013     1     5    1155             1200          -5    1241           1306
## 7  2013     1     6    2125             2129          -4    2224           2224
## 8  2013     1     7    2124             2129          -5    2212           2224
## 9  2013     1     8    2127             2130          -3    2304           2225
## 10 2013     1     9    2126             2129          -3    2217           2224
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

- Does it matter what order you used `filter()` and `arrange()` if you are using both? Yes, in terms of speed/how much work the function has to do because using `filter()` first reduces the number of rows used in `arrange()` but using `arrange()` first means every row has to be sorted, and either way you have to filter through every row. However, in terms of the result, no. For example, let's say I take a vector 1:10 and filter only even numbers, so it is 2, 4, 6, etc., then arrange it descending, I would get it 10, 8, 6, etc. I filtered through 10 numbers but only had to sort 5 numbers. If I take the same vector and first arrange it descending it would be 10, 9, 8, etc. and then when I filter out odd numbers I would get 10, 8, 6, etc. However, I filtered through 10 numbers and had to sort 10 numbers. While the result is the same, filtering before arranging is faster/more efficient.

Let's test this

```
df <- data.frame(x = 1:100000000)

filter_first <- function() {
  df |>
    filter(x %% 2 == 0) |>
    arrange(desc(x))
}

arrange_first <- function() {
  df |>
    arrange(desc(x)) |>
    filter(x %% 2 == 0)
}
```

```
system.time(filter_first())
```

```
##      user  system elapsed
##    1.220    1.278    3.295
```

```
system.time(arrange_first())
```

```
##      user  system elapsed
##    1.228    1.268    3.638
```

## 3.3 Columns

Four important verbs affect columns without changing rows 1. `mutate()` creates new columns derived from existing ones 2. `select()` changes which are present 3. `rename()` changes names 4. `relocate()` changes positions

### 3.3.1 `mutate()`

Can use to compute gain (how much time delayed flight made up in air) and the speed in mph

```
flights |>
  mutate(
    gain = dep_delay - arr_delay,
    speed = distance / air_time * 60
  )
```

```
## # A tibble: 336,776 x 21
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830           819
## 2  2013     1     1     533             529           4     850           830
## 3  2013     1     1     542             540           2     923           850
## 4  2013     1     1     544             545          -1    1004          1022
## 5  2013     1     1     554             600          -6     812           837
## 6  2013     1     1     554             558          -4     740           728
```

```
## 7 2013 1 1 555 600 -5 913 854
## 8 2013 1 1 557 600 -3 709 723
## 9 2013 1 1 557 600 -3 838 846
## 10 2013 1 1 558 600 -2 753 745
## # i 336,766 more rows
## # i 13 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>, gain <dbl>, speed <dbl>
```

mutate() adds new rows by default on right side, can use .before to add to left side instead

```
flights |>
  mutate(
    gain = dep_delay - arr_delay,
    speed = distance / air_time * 60,
    .before = 1
  )
```

```
## # A tibble: 336,776 x 21
##   gain speed year month day dep_time sched_dep_time dep_delay arr_time
##   <dbl> <dbl> <int> <int> <int> <int>           <int>         <dbl>    <int>
## 1   -9  370.  2013     1     1     517             515           2      830
## 2  -16  374.  2013     1     1     533             529           4      850
## 3  -31  408.  2013     1     1     542             540           2      923
## 4   17  517.  2013     1     1     544             545          -1     1004
## 5   19  394.  2013     1     1     554             600          -6      812
## 6  -16  288.  2013     1     1     554             558          -4      740
## 7  -24  404.  2013     1     1     555             600          -5      913
## 8   11  259.  2013     1     1     557             600          -3      709
## 9    5  405.  2013     1     1     557             600          -3      838
## 10 -10  319.  2013     1     1     558             600          -2      753
## # i 336,766 more rows
## # i 12 more variables: sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

. indicates that .before is an argument to the function, not the name of a third variable we are creating Can also use .after and in both can use variable name instead of position

For example can add new variables after day column

```
flights |>
  mutate(
    gain = dep_delay - arr_delay,
    speed = distance / air_time * 60,
    .after = day
  )
```

```
## # A tibble: 336,776 x 21
##   year month day gain speed dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int> <dbl> <dbl> <int>           <int>         <dbl>    <int>
## 1  2013     1     1   -9  370.     517             515           2      830
## 2  2013     1     1  -16  374.     533             529           4      850
```



```
## 3 2013 1 1 -31 408. 542 540 2 923
## 4 2013 1 1 17 517. 544 545 -1 1004
## 5 2013 1 1 19 394. 554 600 -6 812
## 6 2013 1 1 -16 288. 554 558 -4 740
## 7 2013 1 1 -24 404. 555 600 -5 913
## 8 2013 1 1 11 259. 557 600 -3 709
## 9 2013 1 1 5 405. 557 600 -3 838
## 10 2013 1 1 -10 319. 558 600 -2 753
## # i 336,766 more rows
## # i 12 more variables: sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## # flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## # distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Alternatively can control which variables are kept with `.keep`, can use “used” argument which specifies to only keep columns involved or created in `mutate()`

```
flights |>
  mutate(
    gain = dep_delay - arr_delay,
    hours = air_time / 60,
    gain_per_hour = gain / hours,
    .keep = "used"
  )
```

```
## # A tibble: 336,776 x 6
##   dep_delay arr_delay air_time gain hours gain_per_hour
##   <dbl>     <dbl>   <dbl> <dbl> <dbl>         <dbl>
## 1         2        11     227   -9  3.78         -2.38
## 2         4        20     227  -16  3.78         -4.23
## 3         2        33     160  -31  2.67        -11.6
## 4        -1       -18     183   17  3.05          5.57
## 5        -6       -25     116   19  1.93          9.83
## 6        -4        12     150  -16  2.5          -6.4
## 7        -5        19     158  -24  2.63         -9.11
## 8        -3       -14      53   11  0.883         12.5
## 9        -3        -8     140    5  2.33          2.14
## 10       -2         8     138  -10  2.3         -4.35
## # i 336,766 more rows
```

### 3.2.2 select()

Allows you to zoom in on useful subset of variables based on their names

Select by names

```
flights |>
  select(year, month, day)
```

```
## # A tibble: 336,776 x 3
##   year month day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
```

```
## 3 2013 1 1
## 4 2013 1 1
## 5 2013 1 1
## 6 2013 1 1
## 7 2013 1 1
## 8 2013 1 1
## 9 2013 1 1
## 10 2013 1 1
## # i 336,766 more rows
```

Select all columns between year and day (inclusive)

```
flights |>
  select(year:day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # i 336,766 more rows
```

Select all columns except those from year to day (inclusive)

```
flights |>
  select(!year:day)
```

```
## # A tibble: 336,776 x 16
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
##   <int>         <int>         <dbl>   <int>         <int>         <dbl> <chr>
## 1     517           515           2     830           819           11 UA
## 2     533           529           4     850           830           20 UA
## 3     542           540           2     923           850           33 AA
## 4     544           545          -1    1004          1022          -18 B6
## 5     554           600          -6     812           837          -25 DL
## 6     554           558          -4     740           728           12 UA
## 7     555           600          -5     913           854           19 B6
## 8     557           600          -3     709           723          -14 EV
## 9     557           600          -3     838           846           -8 B6
## 10    558           600          -2     753           745            8 AA
## # i 336,766 more rows
## # i 9 more variables: flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Recommend use ! instead of -

Select all columns that are characters

```
flights |>
  select(where(is.character))
```

```
## # A tibble: 336,776 x 4
##   carrier tailnum origin dest
##   <chr>    <chr>   <chr> <chr>
## 1 UA      N14228  EWR   IAH
## 2 UA      N24211  LGA   IAH
## 3 AA      N619AA   JFK   MIA
## 4 B6      N804JB   JFK   BQN
## 5 DL      N668DN   LGA   ATL
## 6 UA      N39463   EWR   ORD
## 7 B6      N516JB   EWR   FLL
## 8 EV      N829AS   LGA   IAD
## 9 B6      N593JB   JFK   MCO
## 10 AA     N3ALAA   LGA   ORD
## # i 336,766 more rows
```

Helper functions you can use within select() - starts\_with("abc") matches names that begin with "abc" - ends\_with("xyz") matches names that end with "xyz" - contains("ijk") matches names that contain "ijk" - num\_range("x", 1:3) matches x1, x2, and x3

For more details

```
?select
```

Can use matches() with regex

Can rename variables as you select() with =, new name on LHS, old on RHS

```
flights |>
  select(tail_num = tailnum)
```

```
## # A tibble: 336,776 x 1
##   tail_num
##   <chr>
## 1 N14228
## 2 N24211
## 3 N619AA
## 4 N804JB
## 5 N668DN
## 6 N39463
## 7 N516JB
## 8 N829AS
## 9 N593JB
## 10 N3ALAA
## # i 336,766 more rows
```

### 3.3.3 rename()

If want to keep all existing variables and rename a few, use `rename()` instead of `select()`

```
flights |>
  rename(tail_num = tailnum)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
## 7  2013     1     1     555             600          -5     913             854
## 8  2013     1     1     557             600          -3     709             723
## 9  2013     1     1     557             600          -3     838             846
## 10 2013     1     1     558             600          -2     753             745
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tail_num <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

For automated column name cleaning check out `janitor::clean_names()`

### 3.3.4 relocate()

By default moves variables to the front

```
flights |>
  relocate(time_hour, air_time)
```

```
## # A tibble: 336,776 x 19
##   time_hour          air_time year month   day dep_time sched_dep_time
##   <dtm>          <dbl> <int> <int> <int>   <int>         <int>
## 1 2013-01-01 05:00:00      227  2013     1     1     517             515
## 2 2013-01-01 05:00:00      227  2013     1     1     533             529
## 3 2013-01-01 05:00:00      160  2013     1     1     542             540
## 4 2013-01-01 05:00:00      183  2013     1     1     544             545
## 5 2013-01-01 06:00:00      116  2013     1     1     554             600
## 6 2013-01-01 05:00:00      150  2013     1     1     554             558
## 7 2013-01-01 06:00:00      158  2013     1     1     555             600
## 8 2013-01-01 06:00:00       53  2013     1     1     557             600
## 9 2013-01-01 06:00:00      140  2013     1     1     557             600
## 10 2013-01-01 06:00:00      138  2013     1     1     558             600
## # i 336,766 more rows
## # i 12 more variables: dep_delay <dbl>, arr_time <int>, sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, distance <dbl>, hour <dbl>, minute <dbl>
```

Can also specify where to put them using `.before` and `.after` like in `mutate()`

```
flights |>
  relocate(year:dep_time, .after = time_hour)
```

```
## # A tibble: 336,776 x 19
##   sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier flight
##   <int>         <dbl>   <int>         <int>         <dbl> <chr>   <int>
## 1         515           2     830           819          11 UA     1545
## 2         529           4     850           830          20 UA     1714
## 3         540           2     923           850          33 AA     1141
## 4         545          -1    1004          1022         -18 B6      725
## 5         600          -6     812           837         -25 DL      461
## 6         558          -4     740           728          12 UA     1696
## 7         600          -5     913           854          19 B6      507
## 8         600          -3     709           723         -14 EV     5708
## 9         600          -3     838           846          -8 B6       79
## 10        600          -2     753           745           8 AA     301
## # i 336,766 more rows
## # i 12 more variables: tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>, year <int>,
## #   month <int>, day <int>, dep_time <int>
```

```
flights|>
  relocate(starts_with("arr"), .before = dep_time)
```

```
## # A tibble: 336,776 x 19
##   year month   day arr_time arr_delay dep_time sched_dep_time dep_delay
##   <int> <int> <int>   <int>     <dbl>   <int>         <int>         <dbl>
## 1  2013     1     1     830         11     517           515           2
## 2  2013     1     1     850         20     533           529           4
## 3  2013     1     1     923         33     542           540           2
## 4  2013     1     1    1004        -18     544           545          -1
## 5  2013     1     1     812        -25     554           600          -6
## 6  2013     1     1     740         12     554           558          -4
## 7  2013     1     1     913         19     555           600          -5
## 8  2013     1     1     709        -14     557           600          -3
## 9  2013     1     1     838         -8     557           600          -3
## 10 2013     1     1     753           8     558           600          -2
## # i 336,766 more rows
## # i 11 more variables: sched_arr_time <int>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

### 3.3.5 Exercises

1. Compare `dep_time`, `sched_dep_time`, and `dep_delay`, how would you expect those 3 numbers to be related? They are all related to the departure time, specifically:  $\text{dep\_delay} = \text{dep\_time} - \text{sched\_dep\_time}$

```
flights |>
  select(contains("dep"))
```

```
## # A tibble: 336,776 x 3
##   dep_time sched_dep_time dep_delay
##   <int>         <int>         <dbl>
## 1      517             515           2
## 2      533             529           4
## 3      542             540           2
## 4      544             545          -1
## 5      554             600          -6
## 6      554             558          -4
## 7      555             600          -5
## 8      557             600          -3
## 9      557             600          -3
## 10     558             600          -2
## # i 336,766 more rows
```

2. Brainstorm as many ways as possible to select dep\_time, dep\_delay, arr\_time, and arr\_delay from flights

```
flights |> # Prefix
  select(starts_with("dep") | starts_with("arr"))
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>    <int>    <dbl>
## 1      517         2      830      11
## 2      533         4      850      20
## 3      542         2      923      33
## 4      544        -1     1004     -18
## 5      554        -6      812     -25
## 6      554        -4      740      12
## 7      555        -5      913      19
## 8      557        -3      709     -14
## 9      557        -3      838      -8
## 10     558        -2      753       8
## # i 336,766 more rows
```

```
flights |> # Regex method 1
  select(matches("\\b[a,d,e,p,r]{3}_[a-z]{4,5}\\b"))
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>    <int>    <dbl>
## 1      517         2      830      11
## 2      533         4      850      20
## 3      542         2      923      33
## 4      544        -1     1004     -18
## 5      554        -6      812     -25
## 6      554        -4      740      12
```

```
## 7      555      -5      913      19
## 8      557      -3      709     -14
## 9      557      -3      838      -8
## 10     558      -2      753       8
## # i 336,766 more rows
```

```
flights |> # Regex method 2
  select(matches("\\b(arr|dep)_(time|delay)\\b"))
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>   <int>     <dbl>
## 1      517         2      830         11
## 2      533         4      850         20
## 3      542         2      923         33
## 4      544        -1     1004        -18
## 5      554        -6      812        -25
## 6      554        -4      740         12
## 7      555        -5      913         19
## 8      557        -3      709        -14
## 9      557        -3      838         -8
## 10     558        -2      753          8
## # i 336,766 more rows
```

3. What happens if you specify the name of the same variable multiple times in a `select()` call? It only selects that variable once

```
flights |>
  select(dep_time, dep_time)
```

```
## # A tibble: 336,776 x 1
##   dep_time
##   <int>
## 1      517
## 2      533
## 3      542
## 4      544
## 5      554
## 6      554
## 7      555
## 8      557
## 9      557
## 10     558
## # i 336,766 more rows
```

4. What does `any_of()` function do? Why might it be helpful in conjunction with this vector?

```
variables <- c("year", "month", "day", "dep_delay", "arr_delay")
```

It doesn't check for missing variables, so you can throw in a column that doesn't exist and you won't get an error unlike `all_of()`

```
flights |>
  select(any_of(variables))
```

```
## # A tibble: 336,776 x 5
##   year month   day dep_delay arr_delay
##   <int> <int> <int>     <dbl>     <dbl>
## 1  2013     1     1         2         11
## 2  2013     1     1         4         20
## 3  2013     1     1         2         33
## 4  2013     1     1        -1        -18
## 5  2013     1     1        -6        -25
## 6  2013     1     1        -4         12
## 7  2013     1     1        -5         19
## 8  2013     1     1        -3        -14
## 9  2013     1     1        -3         -8
## 10 2013     1     1        -2          8
## # i 336,766 more rows
```

```
flights |>
  select(all_of(variables))
```

```
## # A tibble: 336,776 x 5
##   year month   day dep_delay arr_delay
##   <int> <int> <int>     <dbl>     <dbl>
## 1  2013     1     1         2         11
## 2  2013     1     1         4         20
## 3  2013     1     1         2         33
## 4  2013     1     1        -1        -18
## 5  2013     1     1        -6        -25
## 6  2013     1     1        -4         12
## 7  2013     1     1        -5         19
## 8  2013     1     1        -3        -14
## 9  2013     1     1        -3         -8
## 10 2013     1     1        -2          8
## # i 336,766 more rows
```

```
# Add a variable that doesn't exist
variables <- c("year", "month", "day", "dep_delay", "arr_delay", "seconds")
```

```
flights |>
  # select(all_of(variables)) Error, can't subset elements that don't exist
  select(any_of(variables))
```

```
## # A tibble: 336,776 x 5
##   year month   day dep_delay arr_delay
##   <int> <int> <int>     <dbl>     <dbl>
## 1  2013     1     1         2         11
## 2  2013     1     1         4         20
## 3  2013     1     1         2         33
## 4  2013     1     1        -1        -18
## 5  2013     1     1        -6        -25
## 6  2013     1     1        -4         12
```



```
## 7 2013 1 1 -5 19
## 8 2013 1 1 -3 -14
## 9 2013 1 1 -3 -8
## 10 2013 1 1 -2 8
## # i 336,766 more rows
```

5. Does result of running following code surprise you? How do select helpers deal with upper and lower case by default and how can you change that?

```
flights |> select(contains("TIME"))
```

```
## # A tibble: 336,776 x 6
##   dep_time sched_dep_time arr_time sched_arr_time air_time time_hour
##   <int>      <int>      <int>      <int>      <dbl> <dtm>
## 1      517          515      830          819      227 2013-01-01 05:00:00
## 2      533          529      850          830      227 2013-01-01 05:00:00
## 3      542          540      923          850      160 2013-01-01 05:00:00
## 4      544          545     1004         1022      183 2013-01-01 05:00:00
## 5      554          600      812          837      116 2013-01-01 06:00:00
## 6      554          558      740          728      150 2013-01-01 05:00:00
## 7      555          600      913          854      158 2013-01-01 06:00:00
## 8      557          600      709          723       53 2013-01-01 06:00:00
## 9      557          600      838          846      140 2013-01-01 06:00:00
## 10     558          600      753          745      138 2013-01-01 06:00:00
## # i 336,766 more rows
```

Yes as it selects rows that contain “time” This is because select helpers default for ignore.case = TRUE You can change that by running the following code

```
flights |> select(contains("TIME", ignore.case = FALSE))
```

```
## # A tibble: 336,776 x 0
```

Now it selects nothing!

6. Rename air\_time to air\_time\_main to indicate units of measurement and move it to the beginning of the data frame

```
flights |>
  rename(air_time_main = air_time) |>
  relocate(air_time_main, .before = 1)
```

```
## # A tibble: 336,776 x 19
##   air_time_main year month   day dep_time sched_dep_time dep_delay arr_time
##   <dbl> <int> <int> <int>   <int>      <int>      <dbl>      <int>
## 1      227  2013     1     1     517          515         2       830
## 2      227  2013     1     1     533          529         4       850
## 3      160  2013     1     1     542          540         2       923
## 4      183  2013     1     1     544          545        -1      1004
## 5      116  2013     1     1     554          600        -6       812
## 6      150  2013     1     1     554          558        -4       740
```

```
## 7      158 2013    1    1    555      600      -5    913
## 8       53 2013    1    1    557      600      -3    709
## 9      140 2013    1    1    557      600      -3    838
## 10     138 2013    1    1    558      600      -2    753
## # i 336,766 more rows
## # i 11 more variables: sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

7. Why doesn't the following work, and what does the error mean?

```
#flights |>
#select(tailnum) |>
#arrange(arr_delay)
```

Error in `arrange(select(flights, tailnum), arr_delay)` : Caused by error: ! object 'arr\_delay' not found  
When you only select tailnum, select outputs a copy of the flights data frame with just the tailnum column, which is the input to the arrange function, which now can't find the arr\_delay column to sort by

### 3.4 The pipe

Real power comes from combining multiple verbs

Find fastest flights from Houston's IAH airport

```
flights |>
  filter(dest == "IAH") |>
  mutate(speed = distance / air_time * 60) |>
  select(year:day, dep_time, carrier, flight, speed) |>
  arrange(desc(speed))
```

```
## # A tibble: 7,198 x 7
##   year month   day dep_time carrier flight speed
##   <int> <int> <int>   <int> <chr>   <int> <dbl>
## 1  2013     7     9       707 UA       226  522.
## 2  2013     8    27      1850 UA      1128  521.
## 3  2013     8    28       902 UA      1711  519.
## 4  2013     8    28      2122 UA      1022  519.
## 5  2013     6    11      1628 UA      1178  515.
## 6  2013     8    27      1017 UA       333  515.
## 7  2013     8    27      1205 UA      1421  515.
## 8  2013     8    27      1758 UA       302  515.
## 9  2013     9    27       521 UA       252  515.
## 10 2013     8    28       625 UA       559  515.
## # i 7,188 more rows
```

Pipe makes this code very easy to read

If we didn't have the pipe... could nest function calls

```

arrange(
  select(
    mutate(
      filter(
        flights,
        dest == "IAH"
      ),
      speed = distance / air_time * 60
    ),
    year:day, dep_time, carrier, flight, speed
  ),
  desc(speed)
)

```

```

## # A tibble: 7,198 x 7
##   year month   day dep_time carrier flight speed
##   <int> <int> <int>   <int> <chr>    <int> <dbl>
## 1  2013     7     9     707 UA        226  522.
## 2  2013     8    27    1850 UA        1128  521.
## 3  2013     8    28     902 UA        1711  519.
## 4  2013     8    28    2122 UA        1022  519.
## 5  2013     6    11    1628 UA        1178  515.
## 6  2013     8    27    1017 UA         333  515.
## 7  2013     8    27    1205 UA        1421  515.
## 8  2013     8    27    1758 UA         302  515.
## 9  2013     9    27     521 UA         252  515.
## 10 2013     8    28     625 UA         559  515.
## # i 7,188 more rows

```

Or use bunch of intermediate objects (the pandas way)

```

flights1 <- filter(flights, dest == "IAH")
flights2 <- mutate(flights1, speed = distance / air_time * 60)
flights3 <- select(flights2, year:day, dep_time, carrier, flight, speed)
arrange(flights3, desc(speed))

```

```

## # A tibble: 7,198 x 7
##   year month   day dep_time carrier flight speed
##   <int> <int> <int>   <int> <chr>    <int> <dbl>
## 1  2013     7     9     707 UA        226  522.
## 2  2013     8    27    1850 UA        1128  521.
## 3  2013     8    28     902 UA        1711  519.
## 4  2013     8    28    2122 UA        1022  519.
## 5  2013     6    11    1628 UA        1178  515.
## 6  2013     8    27    1017 UA         333  515.
## 7  2013     8    27    1205 UA        1421  515.
## 8  2013     8    27    1758 UA         302  515.
## 9  2013     9    27     521 UA         252  515.
## 10 2013     8    28     625 UA         559  515.
## # i 7,188 more rows

```

Pipe code is easier to write/read Shortcut is Cmd + Shift + M

Default is magrittr %>%, part of tidyverse

```
mtcars %>%
  group_by(cyl) %>%
  summarize(n = n())
```

```
## # A tibble: 3 x 2
##   cyl     n
##   <dbl> <int>
## 1     4    11
## 2     6     7
## 3     8    14
```

For simple cases, behaves identical to base pipe `|>`, but base pipe is part of base R and is simpler

## 3.5 Groups

So far worked with rows and columns, dplyr even more powerful when work with groups

### 3.5.1 group\_by()

Use to divide dataset into groups meaningful for analysis

```
flights |>
  group_by(month)
```

```
## # A tibble: 336,776 x 19
## # Groups:   month [12]
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>    <int>         <int>
## 1  2013     1     1     517           515           2      830           819
## 2  2013     1     1     533           529           4      850           830
## 3  2013     1     1     542           540           2      923           850
## 4  2013     1     1     544           545          -1     1004          1022
## 5  2013     1     1     554           600          -6      812           837
## 6  2013     1     1     554           558          -4      740           728
## 7  2013     1     1     555           600          -5      913           854
## 8  2013     1     1     557           600          -3      709           723
## 9  2013     1     1     557           600          -3      838           846
## 10 2013     1     1     558           600          -2      753           745
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Doesn't change data but output shows Groups: month [12], which means subsequent operations will work "by month" `group_by()` adds this class to data frame, which changes behavior of subsequent verbs

### 3.5.2 summarize()

Most important grouped operation is a summary, which if used to calculate one summary statistic reduces each group to one row

Compute average departure delay by month

```
flights |>
  group_by(month) |>
  summarize(
    avg_delay = mean(dep_delay)
  )
```

```
## # A tibble: 12 x 2
##   month avg_delay
##   <int>   <dbl>
## 1     1      NA
## 2     2      NA
## 3     3      NA
## 4     4      NA
## 5     5      NA
## 6     6      NA
## 7     7      NA
## 8     8      NA
## 9     9      NA
## 10    10      NA
## 11    11      NA
## 12    12      NA
```

Whoops, everything is NA! This is because each month contained missing data for some flights, for now, tell `mean()` to ignore missing values with `na.rm = TRUE`

```
flights |>
  group_by(month) |>
  summarize(
    avg_delay = mean(dep_delay, na.rm = TRUE)
  )
```

```
## # A tibble: 12 x 2
##   month avg_delay
##   <int>   <dbl>
## 1     1    10.0
## 2     2    10.8
## 3     3    13.2
## 4     4    13.9
## 5     5    13.0
## 6     6    20.8
## 7     7    21.7
## 8     8    12.6
## 9     9     6.72
## 10    10     6.24
## 11    11     5.44
## 12    12    16.6
```

Can create as many summaries as want in a single call, one very useful is `n()` which returns number of rows in each group

```
flights |>
  group_by(month) |>
  summarize(
    avg_delay = mean(dep_delay, na.rm = TRUE),
    n = n()
  )
```

```
## # A tibble: 12 x 3
##   month avg_delay    n
##   <int>   <dbl> <int>
## 1     1    10.0 27004
## 2     2    10.8 24951
## 3     3    13.2 28834
## 4     4    13.9 28330
## 5     5    13.0 28796
## 6     6    20.8 28243
## 7     7    21.7 29425
## 8     8    12.6 29327
## 9     9     6.72 27574
## 10    10     6.24 28889
## 11    11     5.44 27268
## 12    12    16.6 28135
```

### 3.5.3 The slice\_ functions

Five functions allow for extracting specific rows within each group 1. `df |> slice_head(n = 1)` takes first row from each group 2. `df |> slice_tail(n = 1)` takes last row from each group 3. `df |> slice_min(x, n = 1)` takes row with smallest value of column x 4. `df |> slice_max(x, n = 1)` takes row with largest value of column x 5. `df |> slice_sample(x, n = 1)` takes one random row Can vary n for more rows, or use `prop = 0.1` to select 10% of rows in each group

Find flights most delayed upon arrival at each destination

```
flights |>
  group_by(dest) |>
  slice_max(arr_delay, n = 1) |>
  relocate(dest)
```

```
## # A tibble: 108 x 19
## # Groups:   dest [105]
##   dest  year month  day dep_time sched_dep_time dep_delay arr_time
##   <chr> <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 ABQ   2013     7   22    2145           2007         98     132
## 2 ACK   2013     7   23    1139           800        219     1250
## 3 ALB   2013     1   25     123          2000        323     229
## 4 ANC   2013     8   17    1740          1625         75     2042
## 5 ATL   2013     7   22    2257           759        898     121
## 6 AUS   2013     7   10    2056          1505        351     2347
## 7 AVL   2013     8   13    1156           832        204     1417
## 8 BDL   2013     2   21    1728          1316        252     1839
## 9 BGR   2013    12     1    1504          1056        248     1628
## 10 BHM  2013     4   10     25           1900        325     136
## # i 98 more rows
```

```
## # i 11 more variables: sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

There are 105 destinations but get 108 rows... Why? `slice_min()` and `slice_max()` keep tied values, if want exactly one row per group use `with_ties = FALSE`

This is similar to computing max delay with `summarize()` but get whole corresponding row instead of single summary statistic

### 3.5.4 Grouping by multiple variables

Can make a group for each date

```
daily <- flights |>
  group_by(year, month, day)
daily
```

```
## # A tibble: 336,776 x 19
## # Groups:   year, month, day [365]
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>    <int>         <int>
## 1  2013     1     1     517             515           2      830             819
## 2  2013     1     1     533             529           4      850             830
## 3  2013     1     1     542             540           2      923             850
## 4  2013     1     1     544             545          -1     1004            1022
## 5  2013     1     1     554             600          -6      812             837
## 6  2013     1     1     554             558          -4      740             728
## 7  2013     1     1     555             600          -5      913             854
## 8  2013     1     1     557             600          -3      709             723
## 9  2013     1     1     557             600          -3      838             846
## 10 2013     1     1     558             600          -2      753             745
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

When summarize tibble group each summary peels off last group

```
daily_flights <- daily |>
  summarize(n = n())
```

```
## 'summarise()' has grouped output by 'year', 'month'. You can override using the
## '.groups' argument.
```

If happy with this behavior can explicitly request it to suppress message

```
daily_flights <- daily |>
  summarize(
    n = n(),
    .groups = "drop_last"
  )
```

Alternatively can use “drop” to drop all grouping or “keep” to preserve groups

### 3.5.5 Ungrouping

Use `ungroup()` to remove grouping without using `summarize()`

```
daily |>
  ungroup()

## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
## 7  2013     1     1     555             600          -5     913             854
## 8  2013     1     1     557             600          -3     709             723
## 9  2013     1     1     557             600          -3     838             846
##10  2013     1     1     558             600          -2     753             745
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Now try summarizing an ungrouped data frame

```
daily |>
  ungroup() |>
  summarize(
    avg_delay = mean(dep_delay, na.rm = TRUE),
    flights = n()
  )
```

```
## # A tibble: 1 x 2
##   avg_delay flights
##   <dbl>   <int>
## 1    12.6  336776
```

Get one row because `dplyr` treats all rows in ungrouped data frame as one group

### 3.5.6 .by

Can now also use `.by` argument to group within a single function

```
flights |>
  summarize(
    delay = mean(dep_delay, na.rm = TRUE),
    n = n(),
    .by = month
  )
```



```
## # A tibble: 12 x 3
##   month delay    n
##   <int> <dbl> <int>
## 1     1  10.0 27004
## 2    10   6.24 28889
## 3    11   5.44 27268
## 4    12  16.6 28135
## 5     2  10.8 24951
## 6     3  13.2 28834
## 7     4  13.9 28330
## 8     5  13.0 28796
## 9     6  20.8 28243
## 10    7  21.7 29425
## 11    8  12.6 29327
## 12    9   6.72 27574
```

Or if want to group by multiple variables

```
flights |>
  summarize(
    delay = mean(dep_delay, na.rm = TRUE),
    n = n(),
    .by = c(origin, dest)
  )
```

```
## # A tibble: 224 x 4
##   origin dest  delay    n
##   <chr> <chr> <dbl> <int>
## 1 EWR   IAH    11.8  3973
## 2 LGA   IAH     9.06 2951
## 3 JFK   MIA     9.34 3314
## 4 JFK   BQN     6.67  599
## 5 LGA   ATL    11.4 10263
## 6 EWR   ORD    14.6  6100
## 7 EWR   FLL    13.5  3793
## 8 LGA   IAD    16.7  1803
## 9 JFK   MCO    10.6  5464
## 10 LGA  ORD    10.7  8857
## # i 214 more rows
```

.by works with all verbs, has advantage of not needing to use .groups to suppress grouping message or ungroup() when you are done

### 3.5.7 Exercises

1. Which carrier has the worst average delays?

```
flights |>
  summarize(
    delay = mean(dep_delay, na.rm = TRUE),
    .by = carrier
  ) |>
  arrange(desc(delay))
```

```
## # A tibble: 16 x 2
##   carrier delay
##   <chr>   <dbl>
## 1 F9      20.2
## 2 EV      20.0
## 3 YV      19.0
## 4 FL      18.7
## 5 WN      17.7
## 6 9E      16.7
## 7 B6      13.0
## 8 VX      12.9
## 9 OO      12.6
## 10 UA     12.1
## 11 MQ     10.6
## 12 DL      9.26
## 13 AA      8.59
## 14 AS      5.80
## 15 HA      4.90
## 16 US      3.78
```

Challenge: can you disentangle effects of bad airports vs bad carriers?

```
# Hint given by problem
flights |>
  summarize(n(), .by = c(carrier, dest))
```

```
## # A tibble: 314 x 3
##   carrier dest 'n()'
##   <chr>   <chr> <int>
## 1 UA      IAH    6924
## 2 AA      MIA    7234
## 3 B6      BQN     599
## 4 DL      ATL   10571
## 5 UA      ORD    6984
## 6 B6      FLL    6563
## 7 EV      IAD    4048
## 8 B6      MCO    6472
## 9 AA      ORD    6059
## 10 B6     PBI    3161
## # i 304 more rows
```

```
# Get average delays for carrier/airport combo
by_carrier <- flights |>
  summarize(
    carrier_delay = mean(dep_delay, na.rm = TRUE),
    n = n(),
    .by = c(carrier, dest)
  )

# Get average delays for each airport
by_airport <- flights |>
  summarize(
    airport_delay = mean(dep_delay, na.rm = TRUE),
```

```

    .by = dest
  )

# Merge and compare difference between carrier and airport
left_join(by_carrier, by_airport, by = join_by(dest)) |>
  mutate(diff = carrier_delay - airport_delay) |>
  relocate(diff, .after = dest) |>
  arrange(desc(diff))

## # A tibble: 314 x 6
##   carrier dest   diff carrier_delay     n airport_delay
##   <chr>   <chr> <dbl>         <dbl> <int>         <dbl>
## 1 UA     STL    61.5          77.5     2          16.0
## 2 OO     ORD    53.4          67      1          13.6
## 3 OO     DTW    49.2          61      2          11.8
## 4 UA     RDU    47.6          60      1          12.4
## 5 EV     PBI    35.7          48.7     6          13.0
## 6 WN     MSY    19.1          33.4   298          14.2
## 7 9E     BGR    14.5          34      1          19.5
## 8 9E     CLT    14.0          23.2   291           9.22
## 9 EV     CLT    13.9          23.1  2508           9.22
## 10 EV    TYS    13.3          41.8   323          28.5
## # i 304 more rows

```

Using a join we can compare a carrier's performance at each airport to the average delay at an airport, however, as you can see, some of these carriers have relatively few flights to a certain airport, so we can't disentangle in all cases

Let's drop cases with relatively few flights

```

left_join(by_carrier, by_airport, by = join_by(dest)) |>
  mutate(diff = carrier_delay - airport_delay) |>
  relocate(diff, .after = dest) |>
  arrange(desc(diff)) |>
  filter(n >= 10)

## # A tibble: 270 x 6
##   carrier dest   diff carrier_delay     n airport_delay
##   <chr>   <chr> <dbl>         <dbl> <int>         <dbl>
## 1 WN     MSY    19.1          33.4   298          14.2
## 2 9E     CLT    14.0          23.2   291           9.22
## 3 EV     CLT    13.9          23.1  2508           9.22
## 4 EV     TYS    13.3          41.8   323          28.5
## 5 9E     CLE    12.2          25.6   349          13.4
## 6 UA     BQN    11.5          23.9   297          12.4
## 7 9E     DFW    11.0          19.7   379           8.68
## 8 EV     DCA    11.0          21.3  1717          10.3
## 9 9E     ORD     9.97          23.5  1056          13.6
## 10 B6     SLC     9.93          19.0   365           9.03
## # i 260 more rows

```

Here we can confirm WN, EV, and 9E are among some of the worst performing carriers in terms of delays

2. Find the flights that are most delayed upon departure from each destination

```
flights |>
  summarise(delay = mean(dep_delay, na.rm = TRUE), .by = dest) |>
  arrange(desc(delay))
```

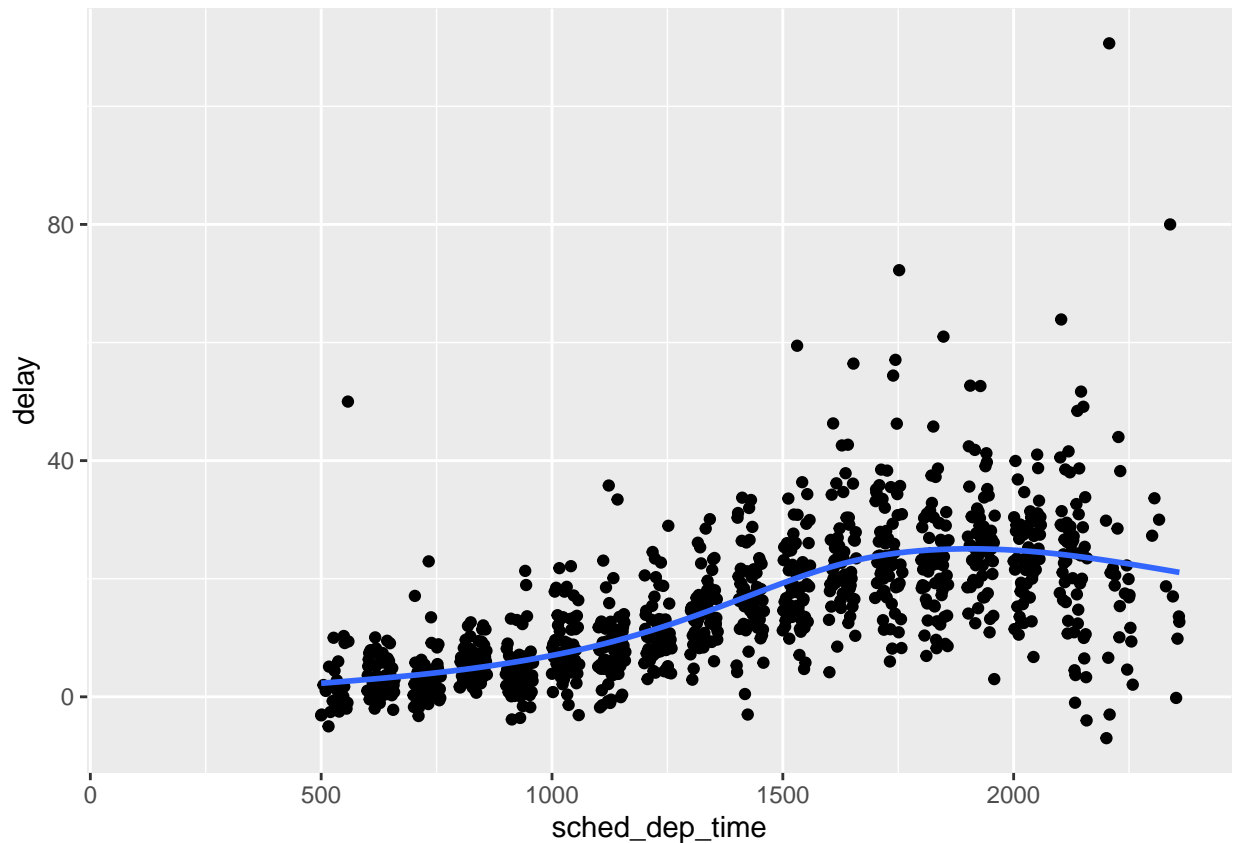
```
## # A tibble: 105 x 2
##   dest    delay
##   <chr> <dbl>
## 1 CAE    35.6
## 2 TUL    34.9
## 3 OKC    30.6
## 4 BHM    29.7
## 5 TYS    28.5
## 6 JAC    26.5
## 7 DSM    26.2
## 8 RIC    23.6
## 9 ALB    23.6
## 10 MSN   23.6
## # i 95 more rows
```

3. How do delays vary over the course of the day? Illustrate with a plot

```
flights |>
  summarise(
    delay = mean(dep_delay, na.rm = TRUE),
    .by = sched_dep_time
  ) |>
  ggplot(
    mapping = aes(
      x = sched_dep_time,
      y = delay
    )
  ) +
  geom_point(na.rm = TRUE) +
  geom_smooth(se = FALSE)
```

```
## 'geom_smooth()' using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```

```
## Warning: Removed 1 row containing non-finite outside the scale range
## ('stat_smooth()').
```



4. What happens if you supply a negative `n` to `slice_min()` and friends? It will grab everything except `n` rows. So `slice_min(-1)` will grab all rows except the smallest row. `slice_max(-1)` will grab all rows except the largest row.

```
flights |>
  group_by(dest) |>
  relocate(dest) |>
  slice_min(dep_delay, n = -1)
```

```
## # A tibble: 336,760 x 19
## # Groups:   dest [103]
##   dest year month day dep_time sched_dep_time dep_delay arr_time
##   <chr> <int> <int> <int> <int>         <int>         <dbl>    <int>
## 1 ABQ   2013   11   20   1948         2000         -12     2236
## 2 ABQ   2013    9   10   1949         2001         -12     2225
## 3 ABQ   2013   11    1   1950         2000         -10     2226
## 4 ABQ   2013   11    5   1950         2000         -10     2310
## 5 ABQ   2013   11   15   1950         2000         -10     2304
## 6 ABQ   2013   12   28   1951         2001         -10     2243
## 7 ABQ   2013    9   18   1951         2001         -10     2210
## 8 ABQ   2013   10   19   1950         1959          -9     2236
## 9 ABQ   2013   10   26   1950         1959          -9     2217
## 10 ABQ  2013   11   10   1951         2000          -9     2258
## # i 336,750 more rows
## # i 11 more variables: sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
```

```
## # flight <int>, tailnum <chr>, origin <chr>, air_time <dbl>, distance <dbl>,
## # hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights |>
  group_by(dest) |>
  relocate(dest) |>
  slice_max(dep_delay, n = -1)
```

```
## # A tibble: 336,762 x 19
## # Groups:   dest [103]
##   dest year month day dep_time sched_dep_time dep_delay arr_time
##   <chr> <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1 ABQ  2013    12   14    2223           2001      142     133
## 2 ABQ  2013    12   17    2220           2001      139     120
## 3 ABQ  2013     7   30    2212           2007      125     57
## 4 ABQ  2013     9    2    2212           2007      125     48
## 5 ABQ  2013     7   23    2206           2007      119    116
## 6 ABQ  2013     5   10    2158           2001      117     53
## 7 ABQ  2013     8    9    2159           2007      112     37
## 8 ABQ  2013     6    8    2148           2001      107      9
## 9 ABQ  2013     9   12    2147           2001      106     25
## 10 ABQ 2013    10   15    2146           2001      105    106
## # i 336,752 more rows
## # i 11 more variables: sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## # flight <int>, tailnum <chr>, origin <chr>, air_time <dbl>, distance <dbl>,
## # hour <dbl>, minute <dbl>, time_hour <dtm>
```

5. Explain what `count()` does in terms of the dplyr verbs, what does `sort` argument to `count()` do?

```
?count
```

It groups by the argument you feed it, and then returns `n()`, the number of rows. If `sort` is `TRUE`, the largest groups will be shown on top.

6. Suppose we have the following tiny data frame:

```
df <- tibble(
  x = 1:5,
  y = c("a", "b", "a", "a", "b"),
  z = c("K", "K", "L", "L", "K")
)
```

a) Write down what you think the output looks like, then check if you are correct and describe what `group_by()` does.

```
df |>
  group_by(y)
```

```
## # A tibble: 5 x 3
## # Groups:   y [2]
##       x y     z
```

```
##      <int> <chr> <chr>
## 1         1 a      K
## 2         2 b      K
## 3         3 a      L
## 4         4 a      L
## 5         5 b      K
```

tibble will show Groups: y [2] because there are two groups, “a” and “b” group\_by() creates groups for each distinct value in y, but it only creates those groups, so the rest of the tibble looks the same

- b) Predict the output, then check if correct, describe what arrange() does, comment how its different from group\_by() in a)

```
df |>
  arrange(y)
```

```
## # A tibble: 5 x 3
##       x y      z
##   <int> <chr> <chr>
## 1     1 a      K
## 2     3 a      L
## 3     4 a      L
## 4     2 b      K
## 5     5 b      K
```

tibble will have re-ordered rows, with rows where x is 1, 3, and 4 appearing first because the tibble is sorted by column y Unlike part (a), the rows are manipulated and no groups are created

- c) Predict output, describe what pipeline does

```
df |>
  group_by(y) |>
  summarize(mean_x = mean(x))
```

```
## # A tibble: 2 x 2
##   y      mean_x
##   <chr>   <dbl>
## 1 a      2.67
## 2 b      3.5
```

y mean\_x a 8/3 b 7/2 Takes the mean of the x values for each group in y, which means taking the mean of all the x values for rows where y is “a” and then doing same for “b”

- d) Do the same then comment on what the message says

```
df |>
  group_by(y, z) |>
  summarize(mean_x = mean(x))
```

```
## 'summarise()' has grouped output by 'y'. You can override using the '.groups'
## argument.
```

```
## # A tibble: 3 x 3
## # Groups:   y [2]
##   y     z     mean_x
##   <chr> <chr> <dbl>
## 1 a     K         1
## 2 a     L        3.5
## 3 b     K        3.5
```

There will be three groups: 1) y = a, z = K 2) y = a, Z = L 3) y = b, z = K Then within each group the mean value of x will be computed y z mean\_x a K 2 a L 4 b K 3.5 After computing the mean, summarise() will return the dataset with the last group (in this case z) dropped, so to override this and keep all groups use .groups = "keep"

e) Do the same, how is output different from d)

```
df |>
  group_by(y, z) |>
  summarize(mean_x = mean(x), .groups = "drop")
```

```
## # A tibble: 3 x 3
##   y     z     mean_x
##   <chr> <chr> <dbl>
## 1 a     K         1
## 2 a     L        3.5
## 3 b     K        3.5
```

The output will be exactly the same as d), except the resulting tibble will have no groups, so if you were to chain another summarize with a pipe, it will calculate the mean for all of the data

```
df |>
  group_by(y, z) |>
  summarize(mean_x = mean(x), .groups = "drop") |>
  summarize(mean_x = mean(mean_x))
```

```
## # A tibble: 1 x 1
##   mean_x
##   <dbl>
## 1  2.67
```

f) Do the same, how are the outputs of the two pipelines different

```
df |>
  group_by(y, z) |>
  summarize(mean_x = mean(x))
```

```
## 'summarise()' has grouped output by 'y'. You can override using the '.groups'
## argument.
```

```
## # A tibble: 3 x 3
## # Groups:   y [2]
##   y     z     mean_x
```



```
##   <chr> <chr>  <dbl>
## 1 a     K      1
## 2 a     L     3.5
## 3 b     K     3.5
```

```
df |>
  group_by(y, z) |>
  mutate(mean_x = mean(x))
```

```
## # A tibble: 5 x 4
## # Groups:   y, z [3]
##       x y     z   mean_x
##   <int> <chr> <chr>   <dbl>
## 1     1 a     K         1
## 2     2 b     K        3.5
## 3     3 a     L        3.5
## 4     4 a     L        3.5
## 5     5 b     K        3.5
```

The first pipeline is the same as d) and it will not have column x in the output, its dimensions will be 3x3  
 The second pipeline creates a new column with the mean x value for each group, so the resulting dimension will be 5x4 since each original row will be retained except it will also have the mean value for its group at the end

### 3.6 Case study: aggregates and sample size

Whenever do any aggregation, good idea to include a count, n(), so you make sure you aren't drawing conclusions from small amounts of data

Demonstrate this with baseball data from Lahman package, specifically will compare batting average (hit / at bat)

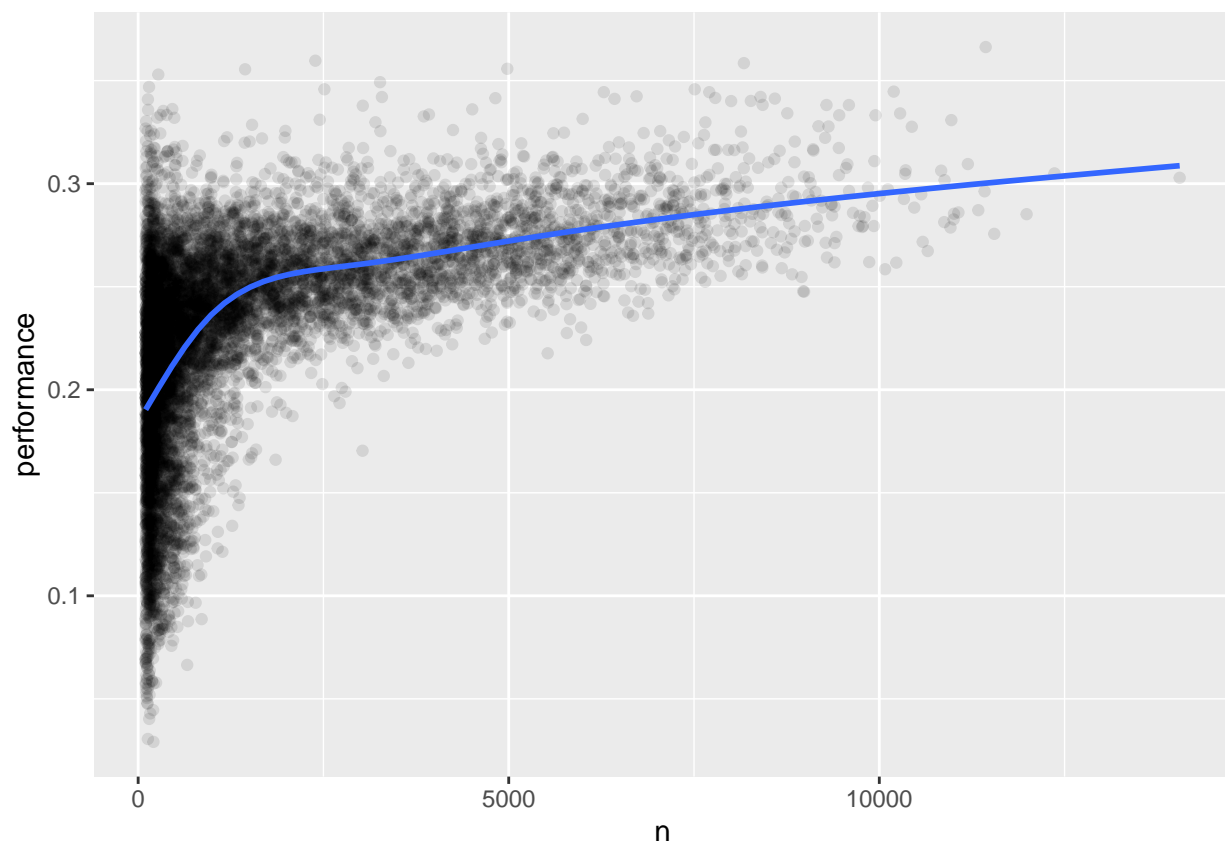
```
batters <- Lahman::Batting |>
  group_by(playerID) |>
  summarize(
    performance = sum(H, na.rm = TRUE) / sum(AB, na.rm = TRUE),
    n = sum(AB, na.rm = TRUE)
  )
batters
```

```
## # A tibble: 20,730 x 3
##   playerID performance      n
##   <chr>         <dbl> <int>
## 1 aardsda01      0         4
## 2 aaronha01    0.305  12364
## 3 aaronto01    0.229   944
## 4 aasedo01      0         5
## 5 abadan01     0.0952    21
## 6 abadfe01     0.111     9
## 7 abadijo01    0.224    49
## 8 abbated01    0.254   3044
## 9 abbeybe01    0.169    225
## 10 abbeych01   0.281   1756
## # i 20,720 more rows
```

When plot batting average (performance) against number of opportunities to hit the ball (n), see two patterns 1. Variation in performance is larger among players with fewer at-bats, shape is very characteristic, whenever you plot a mean (or other summary statistics) vs group size, will see variation decreases as sample size increases (law of large numbers) 2. Positive correlation between skill (performance) and at-bats (n) because teams want to give best batters most opportunities to hit

```
batters |>
  filter(n > 100) |>
  ggplot(aes(x = n, y = performance)) +
  geom_point(alpha = 1 / 10) +
  geom_smooth(se = FALSE)
```

```
## 'geom_smooth()' using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```



Note handy integration between dply and ggplot2, just have to remember to switch from |> to +

This concept also has implications for ranking, if naively sort on desc(performance), will find many players with few at-bats, not necessarily the most skilled players

```
batters |>
  arrange(desc(performance))
```

```
## # A tibble: 20,730 x 3
##   playerID performance      n
##   <chr>          <dbl> <int>
## 1 abramge01            1      1
```

```
## 2 alberan01      1      1
## 3 banisje01      1      1
## 4 bartocl01      1      1
## 5 bassdo01       1      1
## 6 birasst01      1      2
## 7 bruneju01      1      1
## 8 burnscb01      1      1
## 9 cammaer01      1      1
## 10 campsh01      1      1
## # i 20,720 more rows
```

### 3.6.1 Extension

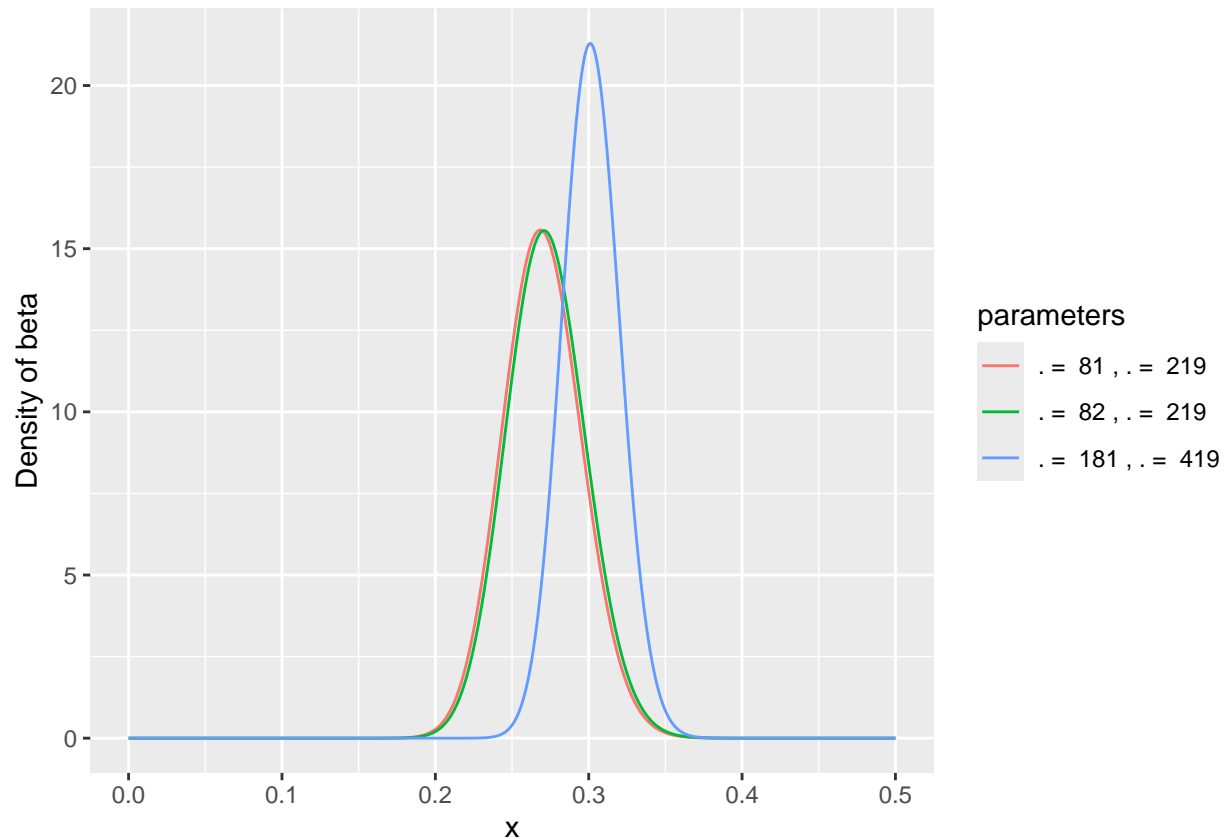
Understanding empirical Bayesian estimation (using baseball statistics)

You can't always throw out data that doesn't meet some minimum, losing information

One approach is beta distribution (probability distribution of probabilities) - If a player's expected pre-season batting average is 0.27 (which will be represented by beta distribution) then bats 100 out of 300, will update beta distribution to be 0.303 less than naive estimate of 0.333 but higher than season-start - Can also see that hitting on first bat (2nd curve) is barely noticable because one hit doesn't tell us much - Helps represent prior expectations and updating based on new evidence

```
# Will plot three curves:
# 1) alpha = 81, beta = 219
# 2) alpha = 82, beta = 219
# 3) alpha = 181, beta = 419
tibble(a = c(81, 82, 81 + 100), b = c(219, 219, 219 + 200)) |>
  rowwise() |> # Go row by row
  mutate(data = list(tibble( # Creates column data with tibbles as value
    # List serves as container to hold nested tibbles in parent table
    # Each parent row values gets broadcasted to the nested rows in its tibble
    x = seq(0, 1, 0.001), # Points at increments 0.001
    y = dbeta(x, a, b), # Apply to the beta distribution function
    parameters = paste("\u03B1 = ", a, ", \u03B2 = ", b) # For label
  ))) |>
  unnest(data) |> # Expands tibbles to data frame
  # Turning parameter into factor will preserve legend order
  mutate(parameters = factor(parameters, levels = unique(parameters))) |>

  # Plot distributions
  ggplot(aes(x, y, color = parameters)) +
  geom_line(na.rm = TRUE) +
  xlim(0, 0.5) +
  ylab("Density of beta")
```



Related method is empirical Bayesian estimation, where beta distribution is used to improve large set of estimates, and as long as you have a lot of examples, you don't need prior expectations

Prepare and clean data first

```
career <- Lahman::Batting |>
  filter(AB > 0) |>
  # Get rid of pitchers (weak batters)
  anti_join(Lahman::Pitching, by = "playerID") |>
  summarize(H = sum(H), AB = sum(AB), .by = playerID) |>
  mutate(average = H / AB)

# Use names as identifier instead
career <- Lahman::People |>
  as_tibble() |>
  select(playerID, nameFirst, nameLast) |>
  unite(name, nameFirst, nameLast, sep = " ") |> # Pastes columns into one
  inner_join(career, by = "playerID") |>
  select(-playerID)

career

## # A tibble: 10,056 x 4
##   name                H    AB average
##   <chr>             <int> <int>   <dbl>
## 1 Hank Aaron       3771 12364  0.305
## 2 Tommie Aaron      216   944   0.229
```

```
## 3 Andy Abad          2    21 0.0952
## 4 John Abadie        11    49 0.224
## 5 Ed Abbaticchio    772  3044 0.254
## 6 Fred Abbott       107   513 0.209
## 7 Jeff Abbott       157   596 0.263
## 8 Kurt Abbott       523  2044 0.256
## 9 Ody Abbott        13    70 0.186
## 10 Frank Abercrombie 0     4 0
## # i 10,046 more rows
```

Who are best players in history? Let's check players with highest batting average...

```
career |>
  arrange(desc(average)) |>
  head(5)
```

```
## # A tibble: 5 x 4
##   name          H    AB average
##   <chr>      <int> <int>   <dbl>
## 1 Jeff Banister    1     1     1
## 2 Doc Bass         1     1     1
## 3 Steve Biras      2     2     1
## 4 C. B. Burns      1     1     1
## 5 Jackie Gallagher 1     1     1
```

Probably just got lucky...

What about the worst?

```
career |>
  arrange(average) |>
  head(5)
```

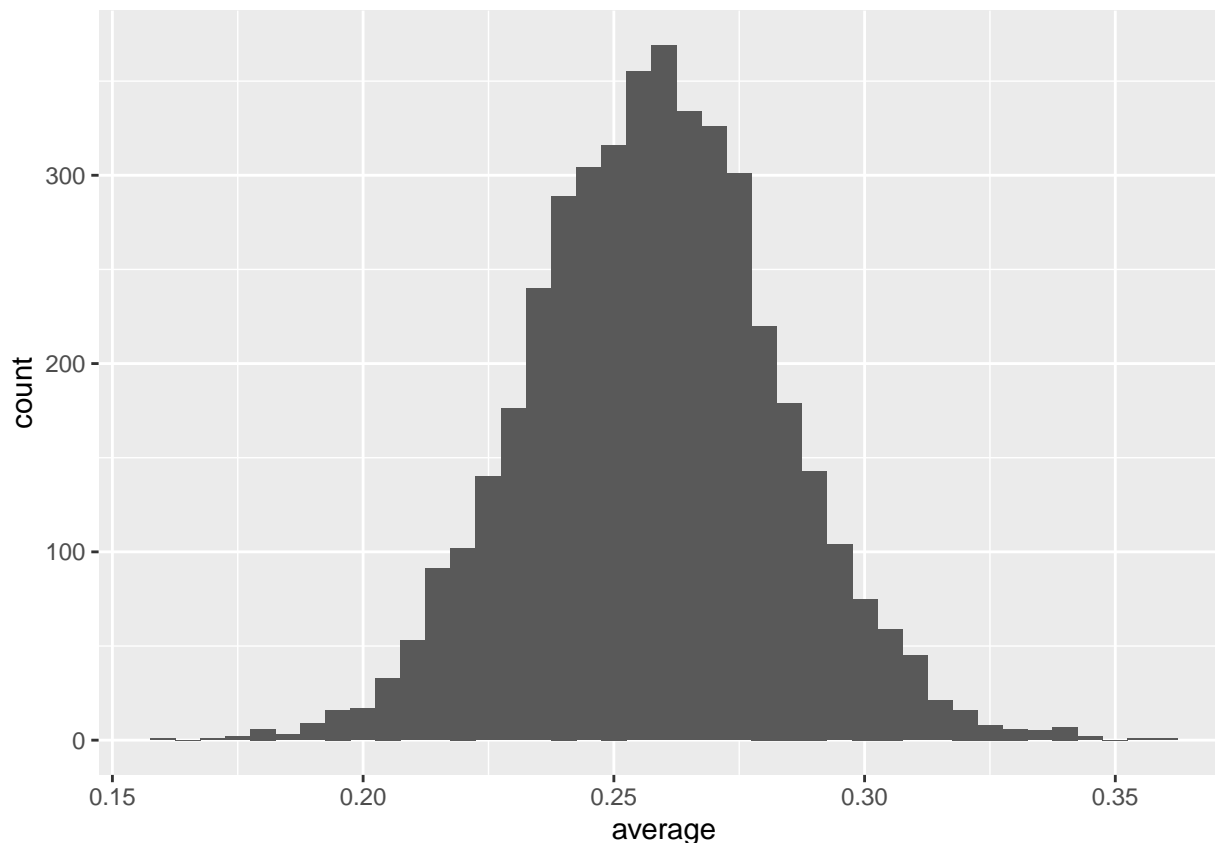
```
## # A tibble: 5 x 4
##   name          H    AB average
##   <chr>      <int> <int>   <dbl>
## 1 Frank Abercrombie 0     4     0
## 2 Horace Allen      0     7     0
## 3 Pete Allen        0     4     0
## 4 Walter Alston     0     1     0
## 5 Trey Amburgey     0     4     0
```

Average here is not a great estimate

Step 1: estimate a prior from all your data

Let's filter out noise (< 500 at bats) and look at the distribution of batting averages across players

```
career |>
  filter(AB >= 500) |>
  ggplot(aes(average)) +
  geom_histogram(binwidth = 0.005)
```



Step 1: estimate a beta prior using this data; estimating from data currently analyzing not typical, usually decide ahead of time Empirical Bayes is an approximation of more exact Bayesian methods, with the amount of data we have (a lot), it's very good

So far data looks good, if it had two or more peaks then might need mixture of betas/more complicated model

Need to pick  $\alpha_0$  and  $\beta_0$ , the “hyper-parameters” of the model

Use `fitdistr` function from MASS to fit probability distribution to data

```
# Filter players we have good estimate of (just like graph)
career_filtered <- career |>
  filter(AB >= 500)

m <- MASS::fitdistr(
  career_filtered$average, # Numeric vector with data
  dbeta, # Function returning a density
  start = list(shape1 = 1, shape2 = 2)
) # List with parameters to be optimized
```

```
## Warning in densfun(x, parm[1], parm[2], ...): NaNs produced
```

```
alpha0 <- m$estimate[1]
beta0 <- m$estimate[2]

alpha0
```

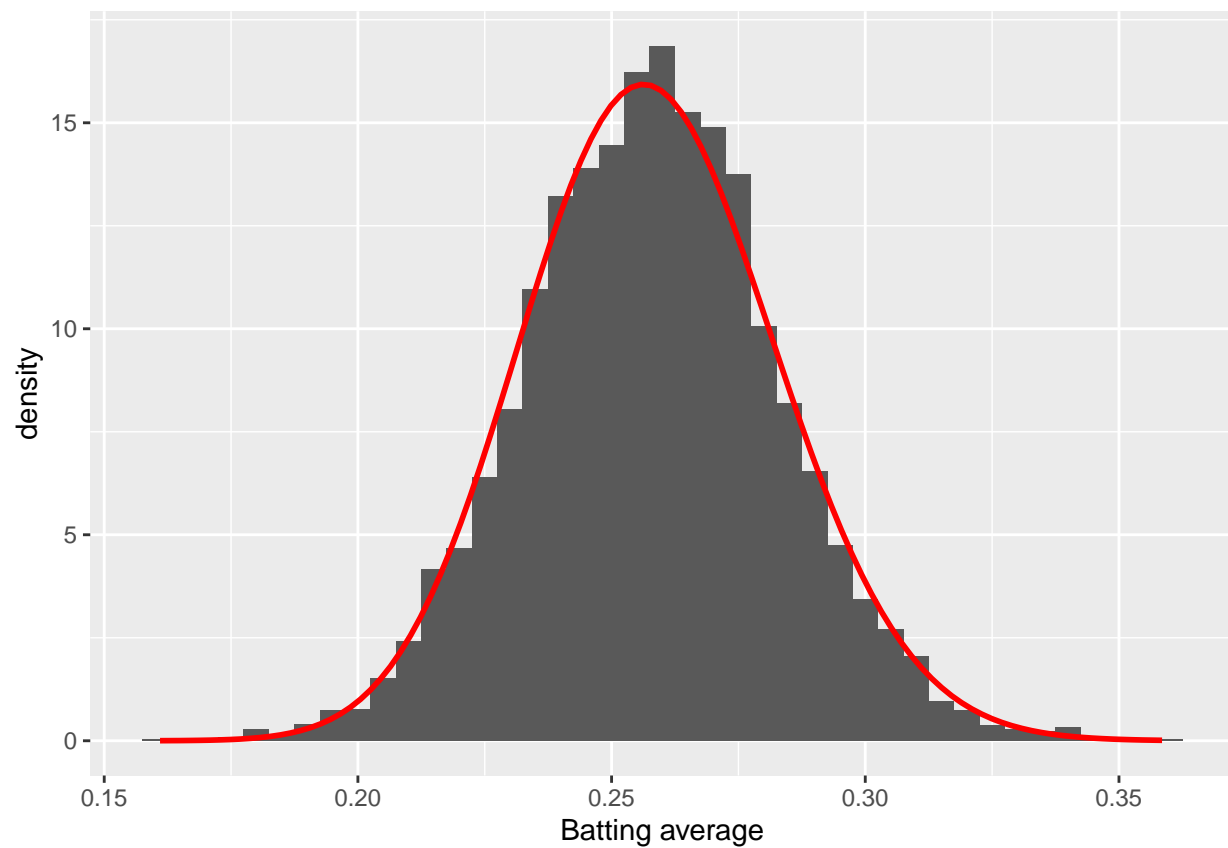
```
## shape1
## 78.59667
```

```
beta0
```

```
## shape2
## 226.1636
```

Lets fit to data

```
ggplot(data = career_filtered) +
  geom_histogram(
    mapping = aes(average, y = after_stat(density)),
    binwidth = 0.005
  ) +
  stat_function(
    fun = function(x) dbeta(x, alpha0, beta0),
    color = "red",
    linewidth = 1
  ) +
  xlab("Batting average")
```



Step 2: use that distribution as prior for each individual estimate  
Start with overall prior and update based on individual evidence

This would yield a higher estimate for the batter who has 300 hits in 1000 AB's than the batter who has 4 hits in 10 AB's, whereas without this method the 4/10 would be considered higher

Perform calculation for all batters

```
career_eb <- career |>
  mutate(eb_estimate = (H + alpha0) / (AB + alpha0 + beta0))
```

Results:

Now we can ask who are the best batters?

```
career_eb |>
  arrange(desc(eb_estimate)) |>
  head(5)
```

```
## # A tibble: 5 x 5
##   name                H    AB average eb_estimate
##   <chr>             <int> <int>   <dbl>      <dbl>
## 1 Rogers Hornsby      2930  8173   0.358      0.355
## 2 Shoeless Joe Jackson 1772  4981   0.356      0.350
## 3 Ed Delahanty        2597  7510   0.346      0.342
## 4 Billy Hamilton      2164  6283   0.344      0.340
## 5 Harry Heilmann      2660  7787   0.342      0.338
```

Who are the worst batters?

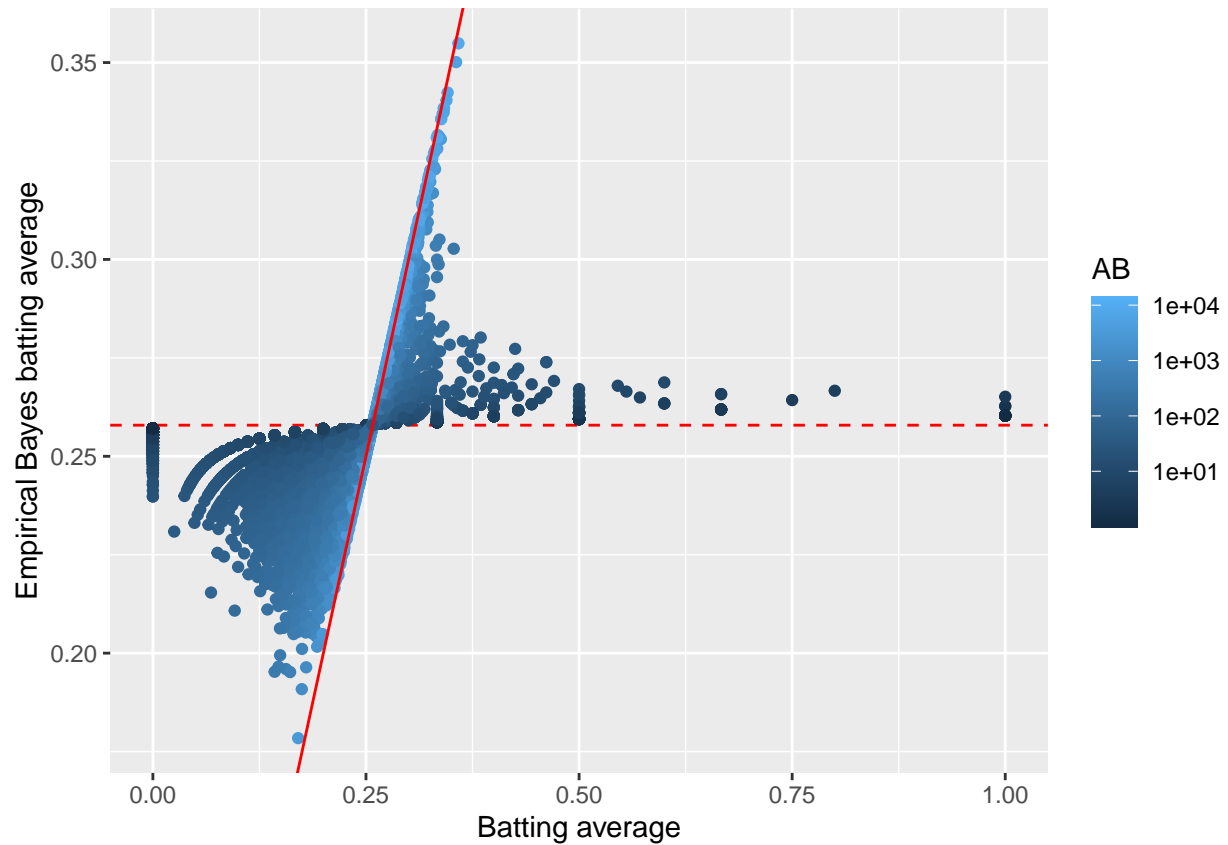
```
career_eb |>
  arrange(eb_estimate) |>
  head(5)
```

```
## # A tibble: 5 x 5
##   name                H    AB average eb_estimate
##   <chr>             <int> <int>   <dbl>      <dbl>
## 1 Bill Bergen         516  3028   0.170      0.178
## 2 Ray Oyler           221  1265   0.175      0.191
## 3 John Vukovich        90   559   0.161      0.195
## 4 John Humphries       52   364   0.143      0.195
## 5 George Baker         74   474   0.156      0.196
```

Empirical Bayes did not simply pick batters with one or two at-bats, instead finding players who bat well/poorly over a long career

```
ggplot(career_eb, aes(x = average, y = eb_estimate, color = AB)) +
  geom_hline(yintercept = alpha0 / (alpha0 + beta0), color = "red", lty = 2) +
  geom_point() +
  geom_abline(color = "red") +
  scale_color_gradient(transform = "log", breaks = 10 ^ (1:5)) +
  xlab("Batting average") +
  ylab("Empirical Bayes batting average")
```





Horizontal dashed line marks  $y = \alpha_0 / (\alpha_0 + \beta_0)$ , what we would guess would be someone's batting average with no evidence at all. Points above it move down towards it, points below move up towards it. Diagonal line is  $x = y$ , points near it didn't get shrunk by empirical Bayes, these are also the ones with the highest at-bats, light blue; there was enough evidence to believe the naive estimate. This process is sometimes called shrinkage because moved all estimates towards the average, the less evidence the more movement. Extraordinary outliers require extraordinary evidence!